

Methods survey - Modern methods

Dimensionality reduction

Clustering and segmentation

Statistical models

Latent variable models

Dynamics/state space models

Mechanistic models

Inferring connectivity

Spatiotemporal pattern analysis

Machine learning models

Comparing data to ML models

Variational inference/ELBO

Learning and inference techniques

Information theoretic methods

Dimensionality reduction (I)


Project or embed data onto “lower-dimensional” subspace or manifold

Often data = population activity patterns, e.g. firing rates estimated by binning spike trains

There are MANY definitions of dimensionality

Three important classes

- Ambient dimensionality: How many numbers per data point in experimental recording
 - Usually determined by experiment setup rather than intrinsic structure in data
- Linear dimensionality: Number of basis vectors needed to span subspace containing data — captures linear structure in data
 - PCA-dimensionality: # of principal components needed to explain X% of total variance
 - Dependent on choice of X
 - Participation ratio: Continuous-valued measure of linear dimensionality
 - [Gao et al bioRxiv 2017](#), [Recanatesi et al Patterns 2022](#)
- Nonlinear dimensionality: How many numbers are needed to specify where a data point lives in the ambient space — captures nonlinear structure in data
 - Many methods to estimate this from data
 - Often called “intrinsic” dimensionality



Dimensionality reduction (II) - Linear methods

- Principal Components Analysis (PCA)
 - Compute covariance matrix of dataset matrix
 - Perform eigendecomposition
 - Returns list of (orthogonal) eigenvectors sorted by how much variance in the data they explain
 - Project data onto top K eigenvectors to reduce to K dimensions
- Singular Value Decomposition (SVD)
 - Mathematically equivalent to PCA applied to both rows and columns of dataset
- jPCA ([Churchland et al, Nature 2012](#))
 - PCA variation for identifying rotational dynamics
 - Requires time-varying data
 - Assumes time derivative is rotational matrix (with purely imaginary eigenvalues) applied to system state
- Canonical Correlation Analysis (CCA)
 - Finds projections of dataset A and dataset B that are maximally correlated
 - [Hotelling, Biometrika 1936](#)
 - [Hardoon et al, Neur Comput 2004](#)

Dimensionality reduction (III) - Nonlinear methods

- Multi-dimensional scaling (MDS)
 - [Kruskal, Psychometrika 1964](#)
- Papers using MDS in neuroscience:
 - [Shepard, Science 1987](#)
- Local linear embedding
 - [Roweis & Saul, Science 2000](#)
- IsoMAP
 - [Tenenbaum et al, Science 2000](#)
- t-SNE
 - [van der Maaten & Hinton, JMLR 2008](#)
- Papers using t-SNE
 - [Berman et al, J R Soc Interface 2014](#)
- UMAP
 - [McInnes et al, arXiv 2018](#)
- Gaussian Process Factor Analysis
 - [Yu et al, J. Neurophys. 2009](#)
- LFADS
 - [Pandinarath et al Nature Methods 2018](#)
- Papers using LFADS:
 - [Ames & Churchland, eLife 2019](#)
- Manifold inference of neural dynamics (MIND)
 - [Nieh, Schottdorf et al, Nature 2021](#)
- CEBRA
 - [Schneider et al Nature Methods 2023](#)

Clustering and segmentation

Goal:

- Split time-series into underlying “states”, each corresponding to one time of pattern
 - E.g. an animal locomotion signal might be split into “walking”, “still”
- *True* number of states is usually not very well defined
 - Scientist must usually pick a number or resolution of states, often controlled by model parameter
- Can be done with simple methods such as threshold crossing (e.g. “walking” $\rightarrow$ velocity $> v_{\min}$)
- These days more common to identify “hidden states” via statistical model
 - Capture variability of observations via generative model
 - Well-posed to incorporate extensions, “stickiness” of states, priors over state distributions, etc.

Example papers where clustering/segmentation is used to process neural and/or behavior time-series

- Kemere et al., Journal of Neurophysiology, 2008, <https://doi.org/10.1152/jn.01280.2007>
- Berman et al., Nature Methods, 2014, <https://doi.org/10.1038/nmeth.2975>
- Wilschko et al., Neuron, 2015, <https://doi.org/10.1016/j.neuron.2015.11.031>
- Berman et al., PNAS, 2016, <https://doi.org/10.1073/pnas.1607601113>
- Baldassano et al., Neuron, 2017, <https://doi.org/10.1016/j.neuron.2017.06.041>
- Vidaurre et al., PNAS, 2018, <https://doi.org/10.1073/pnas.1705120115>
- Calhoun, Pillow, Murthy 2019, <https://www.nature.com/articles/s41593-019-0533-x>
- Wilschko et al., Nature Neuroscience, 2020, <https://doi.org/10.1038/s41593-020-00706-3>
- Recanatesi et al., Neuron, 2022, <https://doi.org/10.1016/j.neuron.2021.10.011>
- Markowitz et al., Nature, 2023, <https://doi.org/10.1038/s41586-022-05611-2>

Statistical models

Probabilistic models of a dataset $\mathcal{D}$ i.e.: $\mathcal{D} \sim P(\mathcal{D} | \dots)$

(Random processes perspective; see Lec 2: Fundamentals)

Common uses

Inference of hidden variables
Learning of hidden parameters
Predicting unseen/future observations

Common conventions

$\mathcal{D} \sim P(\mathcal{D} | \Theta)$ Data distribution specified by parameters Θ

$\mathcal{D} \sim P(\mathcal{D} | \mathbf{x}, \Theta)$ Data distribution specified by inputs $\mathbf{x}$ and parameters Θ

$\mathcal{D} \sim P(\mathcal{D} | \mathbf{x}, \Theta, \mathcal{H})$ Data distribution specified by inputs $\mathbf{x}$, parameters Θ , and hyperparameters $\mathcal{H}$

Important time-series models for neuroscience

Homogeneous Poisson process: Standard model of spike times given constant firing rate

Inhomogeneous Poisson process: Standard model of spike times given dynamic firing rate

Linear-nonlinear Poisson (LNP) model: Standard model of stimulus-driven spike train

Generalized linear model (GLM): Slightly more general than LNP, since can produce e.g. continuous outputs

Markov Chain (MC): Memoryless model of probabilistic switching among discrete states

Hidden Markov Model (HMM): Hidden variables follow Markov chain, observations are conditioned on hidden variables

Kalman Filter: A continuous version of an HMM, often used for tracking and inference

Variations on Hidden Markov Models: E.g. observations conditioned on hidden variables AND past observations

Directed graphical models

Pictorial summary of joint probability distribution underlying statistical model.

Very common way to illustrate probability model.

For DAG (directed acyclic graphs), specifies how to sample from the distribution.

NOT the same as computation graphs, although can be related.


Common convention:

$\mathbf{z}_t \rightarrow$ discrete latent


$\mathbf{x}_t \rightarrow$ continuous latent

$\mathbf{y}_t \rightarrow$ observed

Example (Hidden Markov Model)



Another example



Latent variable models

Discrete: $\{\mathbf{y}\} \sim (\{\mathbf{y}\}|\{\mathbf{z}\}, \Theta, \mathcal{H})$

Example: hidden state $\mathbf{z}$ represents discrete animal state (engaged vs not engaged)

Continuous: $\{\mathbf{y}\} \sim (\{\mathbf{y}\}|\{\mathbf{x}\}, \Theta, \mathcal{H})$

Example: hidden state $\mathbf{x}$ represents continuous animal state (e.g. arousal level)

Example papers

Smith & Brown, Neural Computation, 2003, <https://doi.org/10.1162/089976603765202622>

Yu et al., Journal of Neurophysiology, 2009, <https://doi.org/10.1152/jn.90941.2008>

Pillow et al., Neural Computation, 2010, https://doi.org/10.1162/NECO_a_00058

Fox et al., Annals of Applied Statistics, 2011, <https://doi.org/10.1214/10-AOAS395>

Linderman et al., ICML, 2017, <https://doi.org/10.48550/arXiv.1610.08466>


Pandarínath et al., Nature Methods, 2018, <https://doi.org/10.1038/s41592-018-0109-9>

Ashwood et al 2022 <https://www.nature.com/articles/s41593-021-01007-z>

Dynamics and state space models

Model data as evolving according to a dynamical system

- Can apply to either observed variable, hidden variables, or both




Mechanistic models

Inferring connectivity

Typically infer “functional connectivity”

Many different definitions of functional connectivity

- Correlations between signals from different brain
 - Simple, but no “directionality”
- Granger causality
- Mutual information
- Transfer entropy
- Linear dynamical system connectivity
- Instrumental variable approach



Spatiotemporal analyses

Artificial neural networks

Variational inference & ELBO

Technique for computing posterior over latent variable $\mathbf{z}$ given data (observations) $\mathbf{x}$.

Key idea: $p(\mathbf{z}|\mathbf{x})$ is hard to compute. Instead approximate $p(\mathbf{z}|\mathbf{x})$ with $q_\theta(\mathbf{z})$, where q is from family of distributions parameterized by θ .

$p(\mathbf{z})$: prior over latents (easy to compute given $\mathbf{z}$)

$p(\mathbf{x}|\mathbf{z})$: likelihood (easy to compute given $\mathbf{x}$ and $\mathbf{z}$)

$p(\mathbf{x}, \mathbf{z}) = p(\mathbf{x}|\mathbf{z})p(\mathbf{z})$: joint (easy to compute given $\mathbf{x}$ and $\mathbf{z}$)

Bayes' Rule

$$p(\mathbf{z}|\mathbf{x}) = \frac{p(\mathbf{z})p(\mathbf{x}|\mathbf{z})}{p(\mathbf{x})}$$

Hard to compute

Given data $\mathbf{x}$ and family of functions $q_\theta(\mathbf{x})$, find θ that minimizes $\text{DKL}[q_\theta(\mathbf{x}) || p(\mathbf{z}|\mathbf{x})]$.

$$\text{DKL}[q_\theta(\mathbf{x}) || p(\mathbf{z}|\mathbf{x})] = \underbrace{E_q[\log q_\theta(\mathbf{z})] - E_q[p(\mathbf{x}, \mathbf{z})]}_{\text{-ELBO}(\mathbf{z})} + E_q[\log p(\mathbf{x})]$$

Hard to compute

-ELBO($\mathbf{z}$)

To minimize DKL, maximize the
ELBO: $E_q[p(\mathbf{x}, \mathbf{z})] - E_q[\log q_\theta(\mathbf{z})]$

[Blei et al 2016. Variational Inference: A review for statisticians.](#)

Variational autoencoders

Methods survey - Modern methods - Popular learning/inference techniques

Empirical Bayes

Bayes

Gradient descent

Expectation maximization

Simulation-based inference

Laplace approximation

Forward/backward
algorithms

Methods survey - Modern methods - Information theoretic methods

Max-ent/max-caliber

[Schneidman et al *Nature* 2006](#)

[Meshulam et al *Neuron* 2017](#)

Mutual info/transfer entropy

[Aguera y Arcas, Fairhall, Bialek *Neur Comput* 2003](#)

[Kirst, Timme, Battaglia *Nat Commun* 2016](#)

[Fairhall, Shea-Brown, Barreiro *Curr Opin Neurobiol* 2018](#)

Predictive information

[Bialek, Nemenman, Tishby *Neur Comput* 1999](#)

[Palmer et al *PNAS* 2015](#)

Maximally informative dimensions

[Sharpee, Rust, Bialek 2003](#)